

Phase 3: Development Part 1



Name: B.Abinaya

Register Number : 312621243001

College Name : Thangavelu Engineering College

Project 3 : Future Sales Prediction

Project 3: Future Sales Prediction

Objective:

The objective is to create a tool that enables the company to optimize inventory management and make informed business decisions based on data-driven sales predictions. In this part, we understand the problem statement and we created a document on what we have understood and we proceeded ahead with solving the problem. The problem is to develop a predictive model that uses historical sales data to forecast future sales for a retail company.

Problem Definition:

The problem is to develop a predictive model that uses historical sales data to forecast future sales for a retail company. This project involves data preprocessing, feature engineering, model selection, training, and evaluation. This model will predict sales on a certain day after being provided with a certain set of inputs.

Code and Explanation :

Utilize a dataset containing historical sales data. Here a csv file is converted to a DataFrame and the pandas object is used. This code will create a DataFrame using the provided data and column names. Remember to replace the placeholder data with your actual dataset.

This dataset seems to be related to advertising expenditures and their impact on sales. Here are the column meanings:

TV: Advertising budget spent on TV ads.

Radio: Advertising budget spent on radio ads.

Newspaper: Advertising budget spent on newspaper ads.

Sales: Sales generated as a result of the advertising campaign.

Here's how you can implement this in Python using pandas:

#Data Source utilize the dataset

```
import pandas as pd
```

```
df=pd.read_csv(r'Sales.csv')
```

```
print(df)
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
..
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

```
[200 rows x 4 columns]
```

You can use `df.head()` to get the first N rows in Pandas DataFrame.

```
# Print the first few rows of the DataFrame
```

```
print(df.head())
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

We calculate and print the summary statistics of the dataset using `df.describe()` function . The `describe()` method returns description of the data in the DataFrame. If the DataFrame contains numerical data, the description contains these information for each column such as count , mean , std , min , 25% , 50% , 75% , max .

```
# Summary statistics
```

```
summary_stats = df.describe()
```

```
print(summary_stats)
```

	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000
50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000

You can use the `isnull()` or `isna()` method of pandas. `DataFrame` and `Series` to check if each element is a missing value or not. `isnull()` is an alias for `isna()` , and both are used interchangeably.value.

#to check any missing values

```
print(df.isnull().sum())
```

```
TV          0
Radio       0
Newspaper   0
Sales       0
dtype: int64
```

The `fillna()` method replaces the `NULL` values with a specified value. The `fillna()` method returns a new `DataFrame` object unless the `inplace` parameter is set to `True` , in that case the `fillna()` method does the replacing in the original `DataFrame` instead.

#if missing values are their then use this code

```
df.fillna(df.mean(), inplace=True)
```

The `drop_duplicates()` method removes duplicate rows. Use the `subset` parameter if only some specified columns should be considered when looking for duplicates.

#to remove duplicate values

```
df = df.drop_duplicates()
```

Label encoding is a technique used in machine learning and data analysis to convert categorical variables into numerical format. It is particularly useful when working with algorithms that require numerical input, as most machine learning models can only operate on numerical data.

```
from sklearn.metrics import accuracy_score  
from sklearn.preprocessing import StandardScaler, LabelEncoder  
from sklearn.preprocessing import StandardScaler  
  
#to convert categorical variables into numerical format  
labelencoder = LabelEncoder()  
df['class']=labelencoder.fit_transform(df['Sales'])  
print(df.tail(10))
```

	TV	Radio	Newspaper	Sales	class
190	39.5	41.1	5.8	10.8	32
191	75.5	10.8	6.0	11.9	39
192	17.2	4.1	31.6	5.9	7
193	166.8	42.0	3.6	19.6	89
194	149.7	35.6	6.0	17.3	74
195	38.2	3.7	13.8	7.6	14
196	94.2	4.9	8.1	14.0	52
197	177.0	9.3	6.4	14.8	56
198	283.6	42.0	66.2	25.5	118
199	232.1	8.6	8.7	18.4	84

Feature engineering involves a set of techniques that enable us to create new features by combining or transforming the existing ones. These techniques help to highlight the most important patterns and relationships in the data. Here the Total spent is added as a feature using the datas of TV , Radio , Newspaper in the data set .

```
#adding the new feature as new column
```

```
df['Total_Spent'] = df['TV'] + df['Radio'] + df['Newspaper']
```

```
print(df)
```

	TV	Radio	Newspaper	Sales	class	Total_Spent
0	230.1	37.8	69.2	22.1	106	337.1
1	44.5	39.3	45.1	10.4	28	128.9
2	17.2	45.9	69.3	12.0	40	132.4
3	151.5	41.3	58.5	16.5	66	251.3
4	180.8	10.8	58.4	17.9	80	250.0
..
195	38.2	3.7	13.8	7.6	14	55.7
196	94.2	4.9	8.1	14.0	52	107.2
197	177.0	9.3	6.4	14.8	56	192.7
198	283.6	42.0	66.2	25.5	118	391.8
199	232.1	8.6	8.7	18.4	84	249.4

```
[200 rows x 6 columns]
```

Feature Scaling or Standardization: It is a step of Data Pre Processing that is applied to independent variables or features of data. It helps to normalize the data within a particular range. Sometimes, it also helps in speeding up the calculations in an algorithm .

```
# Scaling features (optional, but can be important for some models)
```

```
scaler = StandardScaler()
```

```
df[['TV', 'Radio', 'Newspaper']] = scaler.fit_transform(df[['TV', 'Radio',  
'Newspaper']])
```

```
# Now, the data is preprocessed and ready for modelling
```

```
print(df)
```

	TV	Radio	Newspaper	Sales	class	Total_Spent
0	0.969852	0.981522	1.778945	22.1	106	337.1
1	-1.197376	1.082808	0.669579	10.4	28	128.9
2	-1.516155	1.528463	1.783549	12.0	40	132.4
3	0.052050	1.217855	1.286405	16.5	66	251.3
4	0.394182	-0.841614	1.281802	17.9	80	250.0
..
195	-1.270941	-1.321031	-0.771217	7.6	14	55.7
196	-0.617035	-1.240003	-1.033598	14.0	52	107.2
197	0.349810	-0.942899	-1.111852	14.8	56	192.7
198	1.594565	1.265121	1.640850	25.5	118	391.8
199	0.993206	-0.990165	-1.005979	18.4	84	249.4

[200 rows x 6 columns]

A correlation matrix is a table containing correlation coefficients for many variables. Each cell in the table represents the correlation between two variables. The value might range between -1 and 1.

Correlation matrix

```
correlation_matrix = df.corr()
```

```
print(correlation_matrix)
```

	TV	Radio	Newspaper	Sales	class	Total_Spent
TV	1.000000	0.054809	0.056648	0.901208	0.904861	0.945330
Radio	0.054809	1.000000	0.354104	0.349631	0.346624	0.293211
Newspaper	0.056648	0.354104	1.000000	0.157960	0.144898	0.343059
Sales	0.901208	0.349631	0.157960	1.000000	0.994857	0.924917
class	0.904861	0.346624	0.144898	0.994857	1.000000	0.924750
Total_Spent	0.945330	0.293211	0.343059	0.924917	0.924750	1.000000

Model selection is the process of selecting one final machine learning model from among a collection of candidate machine learning models for a training dataset. Here the ARIMA model is selected . An autoregressive integrated moving average, or ARIMA, is a statistical analysis model that uses time series data to either better understand the data set or to predict future trends . Autoregressive Integrated Moving Average (ARIMA) is a commonly-used local statistical algorithm for time-series forecasting .

```
#import libraries
```

```
from statsmodels.tsa.arima.model import ARIMA
```

```
from itertools import product
```

```
import itertools
```

```
p = 1 # Example value
```

```
d = 1 # Example value
```

```
q = 1 # Example value
```

```
# Create the ARIMA model
```

```
model = ARIMA(y, order=(p, d, q))
```

```
# Fit the model to the data
```

```
model_fit = model.fit()
```

```
# Summary of the model
```

```
print(model_fit.summary())
```

```

=====
SARIMAX Results
=====
Dep. Variable:          Sales      No. Observations:          200
Model:                ARIMA(1, 1, 1)  Log Likelihood            -616.270
Date:                 Tue, 17 Oct 2023  AIC                        1238.541
Time:                 19:31:30      BIC                        1248.421
Sample:                0          HQIC                        1242.539
                        - 200
Covariance Type:          opg
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
ar.L1         -0.0125      0.081      -0.154      0.878      -0.171      0.146
ma.L1         -0.9999      3.737      -0.268      0.789      -8.324      6.324
sigma2        27.9129     104.167      0.268      0.789     -176.251     232.077
=====
Ljung-Box (L1) (Q):                0.00  Jarque-Bera (JB):                3.72
Prob(Q):                           0.95  Prob(JB):                  0.16
Heteroskedasticity (H):              1.02  Skew:                      -0.09
Prob(H) (two-sided):                0.95  Kurtosis:                  2.35
=====

Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

```

Train/Test is a method to measure the accuracy of your model. It is called Train/Test because you split the data set into two sets: a training set and a testing set. 80% for training, and 20% for testing. You train the model using the training set. You test the model using the testing set.

```
# Extract 'Sales' column as a time series
```

```
sales_ts = df['Sales']
```

```
# Train an ARIMA model
```

```
order = (1, 1, 1) # ARIMA(p,d,q) parameters (you may need to tune these)
```

```
model = ARIMA(sales_ts, order=order)
```

```
results = model.fit()
```

```
# Print model summary
```

```
print(results.summary())
```

```
# Optionally, you can make forecasts with the trained model
```

```

# Number of steps to forecast

forecast_steps = 10

forecast = results.get_forecast(steps=forecast_steps)

forecast_mean = forecast.predicted_mean

forecast_ci = forecast.conf_int()

# Print the forecasts

print(forecast_mean)

print(forecast_ci)

```

```

=====
SARIMAX Results
=====
Dep. Variable:          Sales      No. Observations:          200
Model:                ARIMA(1, 1, 1)  Log Likelihood            -616.270
Date:                 Tue, 17 Oct 2023  AIC                        1238.541
Time:                 19:37:57         BIC                        1248.421
Sample:               0              HQIC                       1242.539
                        - 200
Covariance Type:      opg
=====

```

	coef	std err	z	P> z	[0.025	0.975]
ar.L1	-0.0125	0.081	-0.154	0.878	-0.171	0.146
ma.L1	-0.9999	3.737	-0.268	0.789	-8.324	6.324
sigma2	27.9129	104.167	0.268	0.789	-176.251	232.077

```

=====
Ljung-Box (L1) (Q):          0.00  Jarque-Bera (JB):          3.72
Prob(Q):                    0.95  Prob(JB):              0.16
Heteroskedasticity (H):      1.02  Skew:                  -0.09
Prob(H) (two-sided):         0.95  Kurtosis:              2.35
=====

```

Currently, the most popular metrics for evaluating time series forecasting models are MAE, RMSE and AIC. To briefly summarize, both MAE and RMSE measures the magnitude of errors in a set of predictions. The major difference between MAE and RMSE is the impact of the large errors.

MAE: absolute error refers to the magnitude of difference between the prediction of an observation and the true value of that observation. MAE takes the average of absolute errors for a group of predictions and observations as a measurement of the magnitude of errors for the entire group.

MSE: Mean Squared Error (MSE) measures the amount of error in a statistical model. Evaluate the mean squared difference between observed and predicted values. If the model has no errors, the MSE is zero. Its value increases as the model error increases.

RMSE: In machine learning, it is extremely helpful to have a single number to judge a model's performance, whether it be during training, cross-validation, or monitoring after deployment. Root mean square error is one of the most widely used measures for this.

Make predictions on the test set

```
predictions = model_fit.forecast(len(test))
```

Calculate MAE, MSE, RMSE

```
mae = mean_absolute_error(test, predictions)
```

```
mse = mean_squared_error(test, predictions)
```

```
rmse = math.sqrt(mse)
```

#Print the values

```
print(f'Mean Absolute Error (MAE): {mae}')
```


```
print(f'Mean Squared Error (MSE): {mse}')
```

```
print(f'Root Mean Squared Error (RMSE): {rmse}')
```

Mean Absolute Error (MAE): 4.279574757401353
Mean Squared Error (MSE): 27.506268780628666
Root Mean Squared Error (RMSE): 5.244641911573055

A scatter plot (aka scatter chart, scatter graph) uses dots to represent values for two different numeric variables. The position of each dot on the horizontal and vertical axis indicates values for an individual data point. Scatter plots are used to observe relationships between variables.

```
import matplotlib.pyplot as plt  
# Scatter plot between 'TV' and 'Sales'  
plt.scatter(df['TV'], df['Sales'])  
plt.xlabel('TV')  
plt.ylabel('Sales')  
plt.title('Scatter plot of TV vs Sales')  
plt.show()
```



A histogram is a graph that shows the frequency of numerical data using rectangles. The height of a rectangle (the vertical axis) represents the distribution frequency of a variable (the amount, or how often that variable appears).

```
# Histogram of 'Sales'
```


```
plt.hist(df['Sales'], bins=20)
```

```
plt.xlabel('Sales')
```

```
plt.ylabel('Frequency')
```

```
plt.title('Histogram of Sales')
```

```
plt.show()
```



Linear Regression is a supervised learning algorithm in machine learning that supports finding the linear correlation among variables. The result or output of the regression problem is a real or continuous value.

```
from sklearn.linear_model import LinearRegression

# Assuming X and y are your features and target variables
X = df[['TV', 'Radio', 'Newspaper']]
y = df['Sales']

# Initialize and train a Linear Regression model
model = LinearRegression()
model.fit(X, y)

# Get coefficients and intercept
coefficients = model.coef_
intercept = model.intercept_
print(f'Coefficients: {coefficients}')
print(f'Intercept: {intercept}')

Coefficients: [4.66270025  1.58465027  0.00729187]
Intercept: 15.130500000000001
```